

A. Create two linux servers, server1 => install and configure kibana and elasticsearch with basic username and password authentication server2 => install and configure metricbeat.

First of all, the elasticsearch was setup as follows:

```
samana@samana:~$ wget -qO - https://artifacts.elastic.co/GPG-KEY-elasticsearch |
sudo apt-key add -
[sudo] password for samana:
OK
```

```
samana@samana:~$ sudo apt-get install apt-transport-https
Reading package lists... Done
Building dependency tree
Reading state information... Done
The following NEW packages will be installed:
  apt-transport-https
0 upgraded, 1 newly installed, 0 to remove and 0 not upgraded.
Need to get 4,680 B of archives.
After this operation, 162 kB of additional disk space will be used.
Get:1 http://np.archive.ubuntu.com/ubuntu focal-updates/universe amd64 apt-transport-https all 2.0.6 [4,680 B]
Fetched 4,680 B in 1s (9,253 B/s)
Selecting previously unselected package apt-transport-https.
(Reading database ... 190955 files and directories currently installed.)
Preparing to unpack .../apt-transport-https_2.0.6_all.deb ...
Unpacking apt-transport-https (2.0.6) ...
Setting up apt-transport-https (2.0.6) ...
```

```
samana@samana:~$ echo "deb https://artifacts.elastic.co/packages/7.x/apt stable
main" | sudo tee /etc/apt/sources.list.d/elastic-7.x.list
deb https://artifacts.elastic.co/packages/7.x/apt stable main
samana@samana:~$
```

```
samana@samana:~$ sudo apt-get update && sudo apt-get install elasticsearch
Hit:1 http://np.archive.ubuntu.com/ubuntu focal InRelease
Hit:2 http://security.ubuntu.com/ubuntu focal-security InRelease
Hit:3 http://np.archive.ubuntu.com/ubuntu focal-updates InRelease
Get:4 https://artifacts.elastic.co/packages/7.x/apt stable InRelease [13.6 kB]
Hit:5 http://np.archive.ubuntu.com/ubuntu focal-backports InRelease
Get:6 https://artifacts.elastic.co/packages/7.x/apt stable/main amd64 Packages [84.6 kB]
Get:7 https://artifacts.elastic.co/packages/7.x/apt stable/main i386 Packages [64.9 kB]
Hit:8 http://download.virtualbox.org/virtualbox/debian hirsute InRelease
Fetched 163 kB in 2s (95.1 kB/s)
Reading package lists... Done
N: Skipping acquire of configured file 'non-free/binary-i386/Packages' as repository 'http://download.virtualbox.org/virtualbox/debian hirsute InRelease' doesn't support architecture 'i386'
```

Now in order to configure elasticsearch properly, we edit the elasticsearch.yml file

```
samana@samana:/etc/elasticsearch$ ls
elasticsearch.keystore  jvm.options.d      roles.yml
elasticsearch.yml       log4j2.properties  users
jvm.options             role_mapping.yml    users_roles
samana@samana:/etc/elasticsearch$ sudo nano elasticsearch.yml
```

We set the discovery type to single-node

```
GNU nano 4.8      elasticsearch.yml      Modified
discovery.type: single-node
# ===== Elasticsearch Configuration =====>
#
# NOTE: Elasticsearch comes with reasonable defaults for most settings.
#       Before you set out to tweak and tune the configuration, make sure>
#       understand what are you trying to accomplish and the consequences.
#
# The primary way of configuring a node is via this file. This template l>
# the most important settings you may want to configure for a production >
#
# Please consult the documentation for further information on configurati>
# https://www.elastic.co/guide/en/elasticsearch/reference/index.html
#
# ----- Cluster ----->
#
# Use a descriptive name for your cluster:
#
cluster.name: elk-learn
#
# ----- Node ----->
#
# Use a descriptive name for the node:
#
```

We set the network host to 0.0.0.0 to make it available from all other systems

```
#
network.host: 0.0.0.0
#
# By default Elasticsearch listens for HTTP traffic on the first free port
# finds starting at 9200. Set a specific HTTP port here:
#
#http.port: 9200
#
# For more information, consult the network module documentation.
#
# ----- Discovery -----
#
# Pass an initial list of hosts to perform discovery when this node is started
# The default list of hosts is ["127.0.0.1", "127.0.0.1"]
#
#discovery.seed_hosts: ["host1", "host2"]
#
# Bootstrap the cluster using an initial set of master-eligible nodes:
#
```

We set the x-pack security headers to allow username and password based authentication.

```
GNU nano 4.8      elasticsearch.yml      Modified
# By default Elasticsearch is only accessible on localhost. Set a different
# address here to expose this node on the network:
#
network.host: 0.0.0.0
#
# By default Elasticsearch listens for HTTP traffic on the first free port
# finds starting at 9200. Set a specific HTTP port here:
#
#http.port: 9200
#
# For more information, consult the network module documentation.
#
# ----- Discovery -----
#
# Pass an initial list of hosts to perform discovery when this node is started
# The default list of hosts is ["127.0.0.1", "127.0.0.1"]
#
#discovery.seed_hosts: ["host1", "host2"]
#
# Bootstrap the cluster using an initial set of master-eligible nodes:
#
#cluster.initial_master_nodes: ["node-1", "node-2"]
#
# For more information, consult the discovery and cluster formation module
#
# ----- Various -----
#
# Require explicit names when deleting indices:
#
#action.destructive_requires_name: true
xpack.security.enabled: true
xpack.security.authc.api_key.enabled: true
```

Now we set up the passwords for all users as follows

```
samana@samana:/usr/share/elasticsearch$ sudo ./bin/elasticsearch-setup-pas
swords interactive
Initiating the setup of passwords for reserved users elastic,apm_system,ki
bana,kibana_system,logstash_system,beats_system,remote_monitoring_user.
You will be prompted to enter passwords as the process progresses.
Please confirm that you would like to continue [y/N]y

Enter password for [elastic]:
Reenter password for [elastic]:
Enter password for [apm_system]:
Reenter password for [apm_system]:
Enter password for [kibana_system]:
Reenter password for [kibana_system]:
Enter password for [logstash_system]:
Reenter password for [logstash_system]:
Enter password for [beats_system]:
Reenter password for [beats_system]:
Enter password for [remote_monitoring_user]:
Reenter password for [remote_monitoring_user]:
Changed password for user [apm_system]
Changed password for user [kibana_system]
Changed password for user [kibana]
Changed password for user [logstash_system]
Changed password for user [beats_system]
Changed password for user [remote_monitoring_user]
Changed password for user [elastic]
```

Now we enable and start the elasticsearch service

```
samana@samana:~$ sudo systemctl daemon-reload
samana@samana:~$ sudo systemctl enable elasticsearch.service
Synchronizing state of elasticsearch.service with SysV service script with /lib/
systemd/systemd-sysv-install.
Executing: /lib/systemd/systemd-sysv-install enable elasticsearch
Created symlink /etc/systemd/system/multi-user.target.wants/elasticsearch.servic
e → /lib/systemd/system/elasticsearch.service.
samana@samana:~$ sudo systemctl start elasticsearch.service
samana@samana:~$
```

We check the status to verify it is running

```
samana@samana:~$ sudo systemctl status elasticsearch.service
● elasticsearch.service - Elasticsearch
   Loaded: loaded (/lib/systemd/system/elasticsearch.service; enabled; vendor
   Active: active (running) since Thu 2021-12-02 17:34:36 +0545; 23s ago
     Docs: https://www.elastic.co
   Main PID: 11276 (java)
    Tasks: 82 (limit: 14131)
   Memory: 5.8G
   CGroup: /system.slice/elasticsearch.service
           └─11276 /usr/share/elasticsearch/jdk/bin/java -Xshare:auto -Des.ne>
           └─11468 /usr/share/elasticsearch/modules/x-pack-ml/platform/linux->

दि सम्बर 02 17:34:02 samana systemd[1]: Starting Elasticsearch...
दि सम्बर 02 17:34:10 samana systemd-entrypoint[11276]: WARNING: A terminally depr>
दि सम्बर 02 17:34:10 samana systemd-entrypoint[11276]: WARNING: System::setSecuri>
दि सम्बर 02 17:34:10 samana systemd-entrypoint[11276]: WARNING: Please consider r>
दि सम्बर 02 17:34:10 samana systemd-entrypoint[11276]: WARNING: System::setSecuri>
दि सम्बर 02 17:34:13 samana systemd-entrypoint[11276]: WARNING: A terminally depr>
दि सम्बर 02 17:34:13 samana systemd-entrypoint[11276]: WARNING: System::setSecuri>
दि सम्बर 02 17:34:13 samana systemd-entrypoint[11276]: WARNING: Please consider r>
दि सम्बर 02 17:34:13 samana systemd-entrypoint[11276]: WARNING: System::setSecuri>
दि सम्बर 02 17:34:36 samana systemd[1]: Started Elasticsearch.
lines 1-21/21 (END)
```

The server where elasticsearch is configured has ip 192.168.0.5

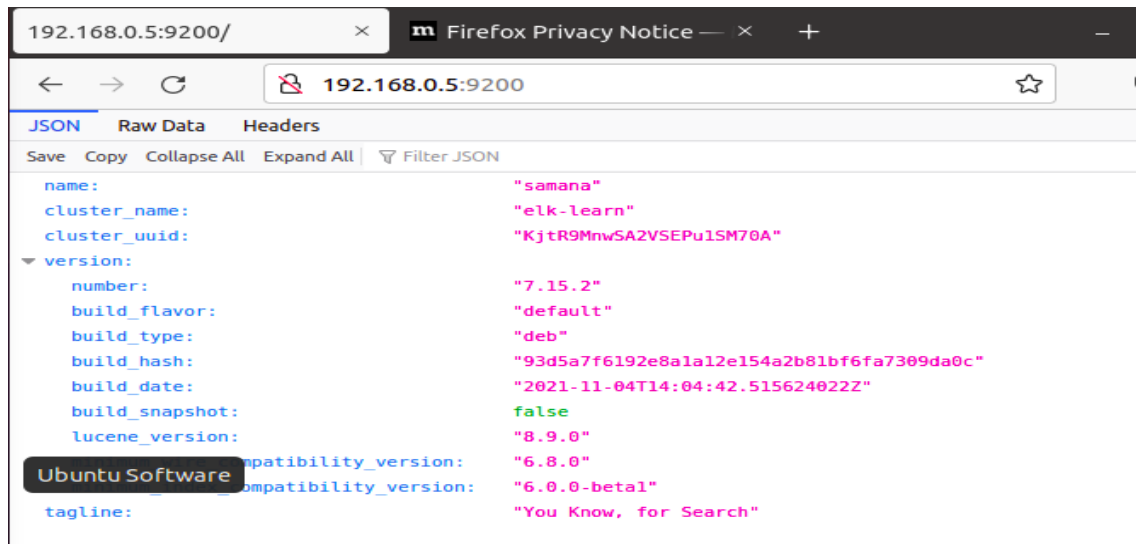
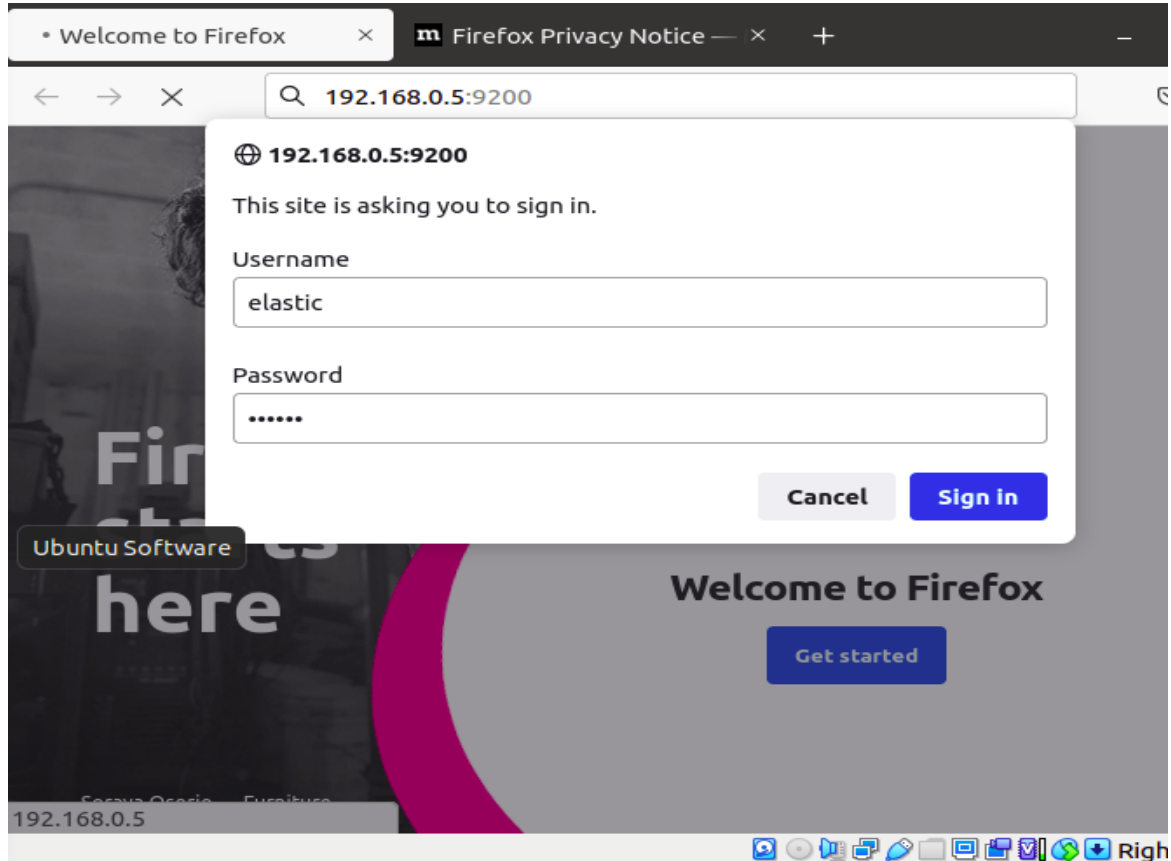
```
samana@samana:~$ ifconfig
enp2s0: flags=4099<UP,BROADCAST,MULTICAST> mtu 1500
    ether 10:7d:1a:3d:19:d2 txqueuelen 1000 (Ethernet)
    RX packets 0 bytes 0 (0.0 B)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 0 bytes 0 (0.0 B)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

lo: flags=73<UP,LOOPBACK,RUNNING> mtu 65536
    inet 127.0.0.1 netmask 255.0.0.0
    inet6 ::1 prefixlen 128 scopeid 0x10<host>
    loop txqueuelen 1000 (Local Loopback)
    RX packets 4889 bytes 476156 (476.1 KB)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 4889 bytes 476156 (476.1 KB)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

wlp1s0: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 1500
    inet 192.168.0.5 netmask 255.255.255.0 broadcast 192.168.0.255
    inet6 fe80::aef4:89f3:7040:1b79 prefixlen 64 scopeid 0x20<link>
    ether a8:6b:ad:1d:9a:09 txqueuelen 1000 (Ethernet)
    RX packets 1104298 bytes 1590594392 (1.5 GB)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 344108 bytes 34328257 (34.3 MB)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

samana@samana:~$ R
```

We access the elasticsearch from our another server (server2)





Now we install kibana which provides interactive GUI for elasticsearch.

```
samana@samana:~$ sudo apt-get install kibana
Reading package lists... Done
Building dependency tree
Reading state information... Done
The following NEW packages will be installed:
  kibana
0 upgraded, 1 newly installed, 0 to remove and 0 not upgraded.
Need to get 288 MB of archives.
After this operation, 786 MB of additional disk space will be used.
Get:1 https://artifacts.elastic.co/packages/7.x/apt stable/main amd64 kibana amd64 7.15.2 [288 MB]
12% [1 kibana 43.2 MB/288 MB 15%] 4,848 kB/s 50s
```

We edit the kibana configuration file and set the server host to anywhere.

```
GNU nano 4.8 kibana.yml Modified
# Kibana is served by a back end server. This setting specifies the port >
#server.port: 5601

# Specifies the address to which the Kibana server will bind. IP address >
# The default is 'localhost', which usually means remote machines will no>
# To allow connections from remote users, set this parameter to a non-loc>
server.host: "0.0.0.0"

# Enables you to specify a path to mount Kibana at if you are running beh>
# Use the `server.rewriteBasePath` setting to tell Kibana if it should re>
# from requests it receives, and to prevent a deprecation warning at star>
# This setting cannot end in a slash.
#server.basePath: ""

# Specifies whether Kibana should rewrite requests that are prefixed with>
# `server.basePath` or require that they are rewritten by your reverse pr>
# This setting was effectively always `false` before Kibana 6.3 and will>
# default to `true` starting in Kibana 7.0.
#server.rewriteBasePath: false
```



We write our previously setup password in our config file ( which should normally be replaced by api keys or encrypted values for security reasons)

```
GNU nano 4.8 kibana.yml Modified
# `server.basePath` or require that they are rewritten by your reverse pr>
# This setting was effectively always `false` before Kibana 6.3 and will
# default to `true` starting in Kibana 7.0.
#server.rewriteBasePath: false

# Specifies the public URL at which Kibana is available for end users. If
# `server.basePath` is configured this URL should end with the same baseP>
#server.publicBaseUrl: ""

# The maximum payload size in bytes for incoming server requests.
#server.maxPayload: 1048576

# The Kibana server's name. This is used for display purposes.
#server.name: "your-hostname"

# The URLs of the Elasticsearch instances to use for all your queries.
#elasticsearch.hosts: ["http://localhost:9200"]

# Kibana uses an index in Elasticsearch to store saved searches, visualiz>
# dashboards. Kibana creates a new index if the index doesn't already exi>
#kibana.index: ".kibana"

# The default application to load.
#kibana.defaultAppId: "home"

# If your Elasticsearch is protected with basic authentication, these set>
# the username and password that the Kibana server uses to perform mainte>
# index at startup. Your Kibana users still need to authenticate with Ela>
# is proxied through the Kibana server.
elasticsearch.username: "kibana system"
elasticsearch.password: "123456"
```

We also setup a encryption key of at least 32 characters

```
GNU nano 4.8 kibana.yml
#logging.dest: stdout

# Set the value of this setting to true to suppress all logging output.
#logging.silent: false

# Set the value of this setting to true to suppress all logging output other than error messages.
#logging.quiet: false

# Set the value of this setting to true to log all events, including system usage information
# and all requests.
#logging.verbose: false

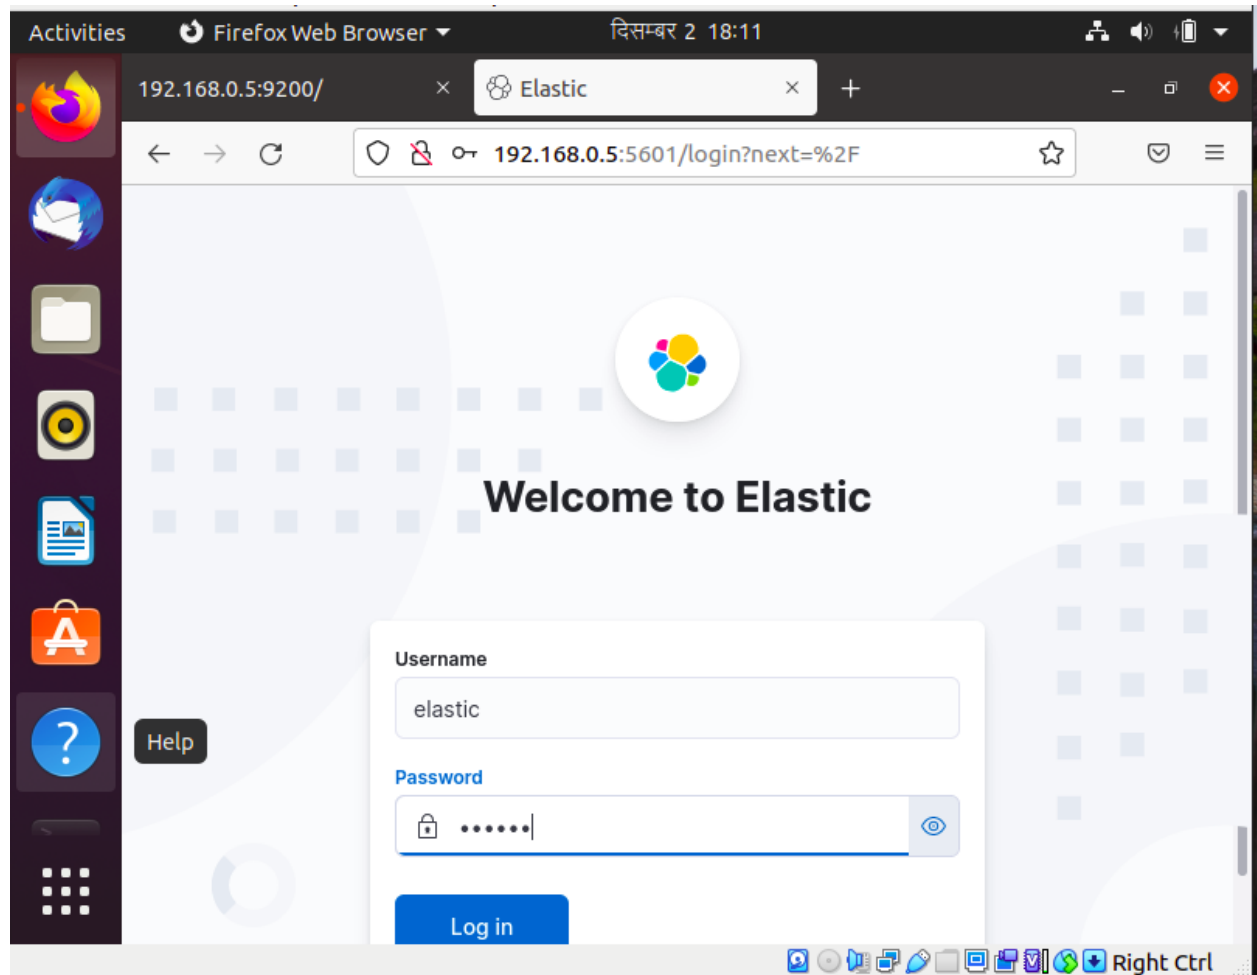
# Set the interval in milliseconds to sample system and process performance
# metrics. Minimum is 100ms. Defaults to 5000.
#ops.interval: 5000

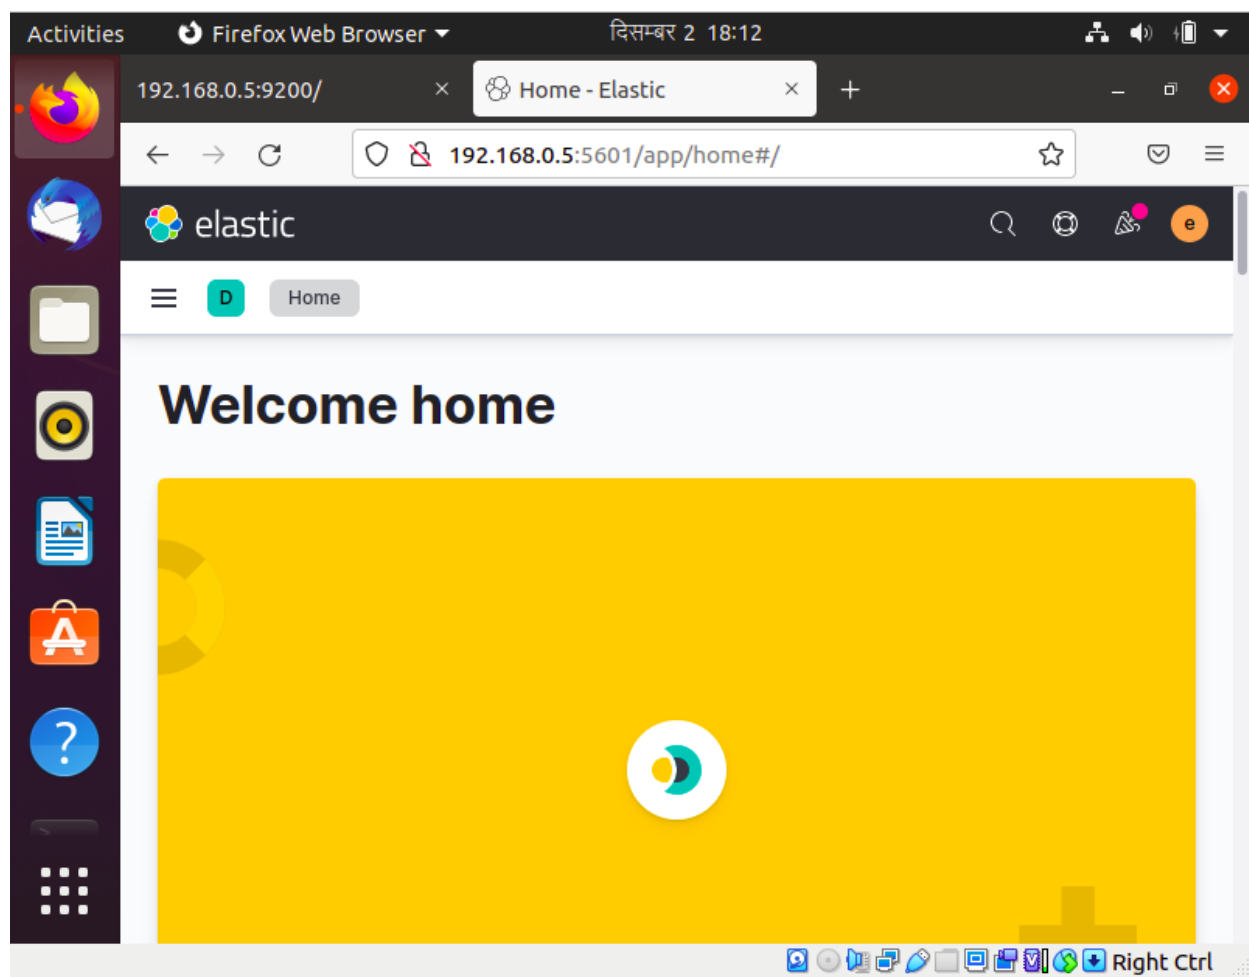
# Specifies locale to be used for all localizable strings, dates and number formats.
# Supported languages are the following: English - en , by default , Chinese - zh-CN .
#i18n.locale: "en"
xpack.encryptedSavedObjects.encryptionKey: "kdjsadsjsjhdshbchshdgdjfdhjdkdebleadsjsjdcharactedddfrfdssdvalue"
```

Now we enable and start the kibana service after editing our config file

```
samana@samana:~$ sudo systemctl enable kibana.service
Synchronizing state of kibana.service with SysV service script with /lib/s
ystemd/systemd-sysv-install.
Executing: /lib/systemd/systemd-sysv-install enable kibana
Created symlink /etc/systemd/system/multi-user.target.wants/kibana.service
→ /etc/systemd/system/kibana.service.
samana@samana:~$ sudo systemctl start kibana.service
samana@samana:~$
```

We can now access our kibana dashboard from our another server





## installing metricbeats in server 2:

```
samana@samana-VM:~$ wget -qO - https://artifacts.elastic.co/GPG-KEY-elasticsearch | sudo apt-key add -
[sudo] password for samana:
OK
samana@samana-VM:~$ sudo apt-get install apt-transport-https
Reading package lists... Done
Building dependency tree
Reading state information... Done
The following NEW packages will be installed:
  apt-transport-https
0 upgraded, 1 newly installed, 0 to remove and 100 not upgraded.
Need to get 1,708 B of archives.
After this operation, 161 kB of additional disk space will be used.
Get:1 http://security.ubuntu.com/ubuntu focal-security/universe amd64 apt-transport-https all 2.0.2ubuntu0.2 [1,708 B]
Fetched 1,708 B in 0s (3,531 B/s)
Selecting previously unselected package apt-transport-https.
(Reading database ... 145618 files and directories currently installed.)
Preparing to unpack .../apt-transport-https_2.0.2ubuntu0.2_all.deb ...
Unpacking apt-transport-https (2.0.2ubuntu0.2) ...
Setting up apt-transport-https (2.0.2ubuntu0.2) ...
samana@samana-VM:~$ echo "deb https://artifacts.elastic.co/packages/7.x/apt stable main" | sudo tee /etc/apt/sources.list.d/elastic-7.x.list
deb https://artifacts.elastic.co/packages/7.x/apt stable main
samana@samana-VM:~$
```

```
samana@samana-VM:~$ sudo apt-get update && sudo apt-get install metricbeat
Get:1 https://artifacts.elastic.co/packages/7.x/apt stable InRelease [13.6 kB]
Get:2 https://artifacts.elastic.co/packages/7.x/apt stable/main i386 Packages [64.9 kB]
Get:3 https://artifacts.elastic.co/packages/7.x/apt stable/main amd64 Packages [84.6 kB]
Err:4 http://security.ubuntu.com/ubuntu focal-security InRelease
  Connection failed [IP: 91.189.88.152 80]
Hit:5 http://np.archive.ubuntu.com/ubuntu focal InRelease
Get:6 http://np.archive.ubuntu.com/ubuntu focal-updates InRelease [114 kB]
Get:7 http://np.archive.ubuntu.com/ubuntu focal-backports InRelease [108 kB]
Get:8 http://np.archive.ubuntu.com/ubuntu focal-updates/main amd64 Packages [1,386 kB]
Ign:9 http://np.archive.ubuntu.com/ubuntu focal-updates/main i386 Packages
Get:10 http://np.archive.ubuntu.com/ubuntu focal-updates/main Translation-en [281 kB]
Get:11 http://np.archive.ubuntu.com/ubuntu focal-updates/main amd64 DEP-11 Metadata [278 kB]
Get:12 http://np.archive.ubuntu.com/ubuntu focal-updates/main DEP-11 48x48 Icons [60.8 kB]
Get:13 http://np.archive.ubuntu.com/ubuntu focal-updates/main DEP-11 64x64 Icons [98.3 kB]
```

We edit the yml file for metricbeat as follows

```
GNU nano 4.8 metricbeat.yml Modif
#cloud.id:

# The cloud.auth setting overwrites the `output.elasticsearch.username` and
# `output.elasticsearch.password` settings. The format is `<user>:<pass>`.
#cloud.auth:

# ===== Outputs =====

# Configure what output to use when sending the data collected by the beat

# ----- Elasticsearch Output -----
output.elasticsearch:
  # Array of hosts to connect to.
  hosts: ["192.168.0.5:9200"]

  # Protocol - either `http` (default) or `https`.
  #protocol: "https"

  # Authentication credentials - either API key or username/password.
  #api_key: "id:api_key"
  username: "elastic"
  password: "123456"
```

We enable and start the metricbeat service.

```
samana@samana-VM:~$ sudo systemctl enable metricbeat
[sudo] password for samana:
Synchronizing state of metricbeat.service with SysV service script with /lib/sy
stemd/systemd-sysv-install.
Executing: /lib/systemd/systemd-sysv-install enable metricbeat
Created symlink /etc/systemd/system/multi-user.target.wants/metricbeat.service
→ /lib/systemd/system/metricbeat.service.
samana@samana-VM:~$ sudo systemctl start metricbeat
samana@samana-VM:~$
```

We setup a metricbeat index which sends the metrics to the kibana dashboard.

```
samana@samana-VM:/etc/metricbeat$ sudo metricbeat setup --template -E 'output.e
lasticsearch.hosts=["192.168.0.5:9200"]'
Flag --template has been deprecated, please use --index-management instead
Overwriting ILM policy is disabled. Set `setup.ilm.overwrite: true` for enablin
g.

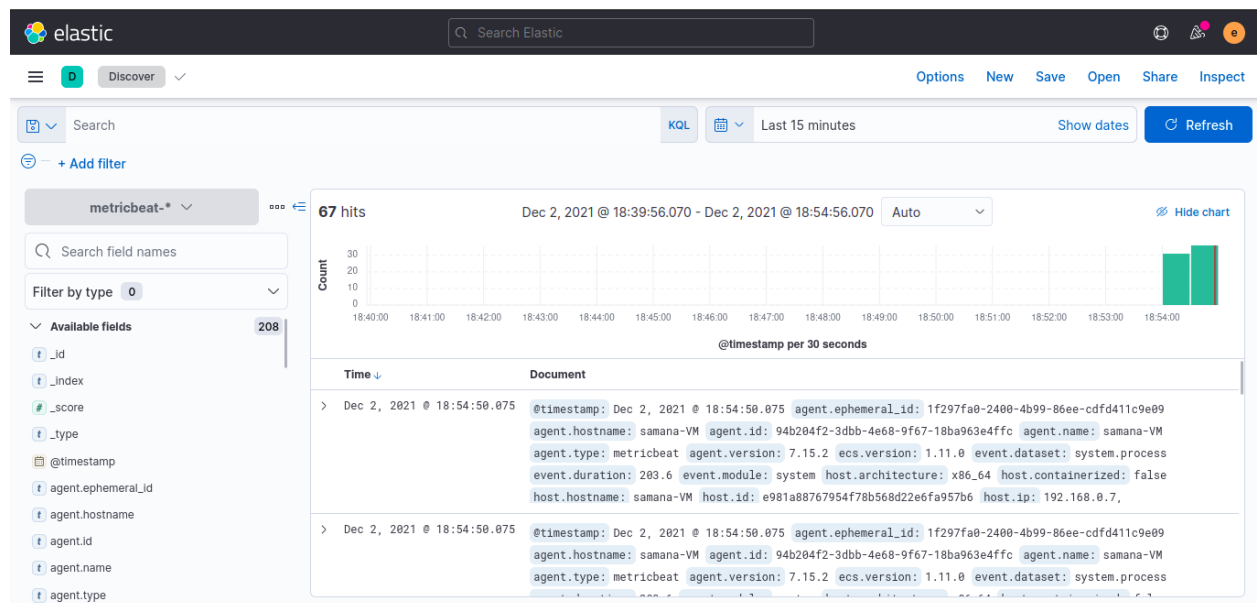
Index setup finished.
samana@samana-VM:/etc/metricbeat$
```

```

samana@samana-VM:/etc/metricbeat$ sudo metricbeat setup -e -E output.elasticsearch.hosts=['192.168.0.5:9200'] -E setup.kibana.host=192.168.0.5:5601
2021-12-02T18:49:01.111+0545 INFO instance/beat.go:665 Home path: [/usr/share/metricbeat] Config path: [/etc/metricbeat] Data path: [/var/lib/metricbeat] Logs path: [/var/log/metricbeat]
2021-12-02T18:49:01.112+0545 INFO instance/beat.go:673 Beat ID: 94b204f2-3dbb-4e68-9f67-18ba963e4ffc
2021-12-02T18:49:01.113+0545 INFO [beat] instance/beat.go:1014 Beat info {"system_info": {"beat": {"path": {"config": "/etc/metricbeat", "data": "/var/lib/metricbeat", "home": "/usr/share/metricbeat", "logs": "/var/log/metricbeat"}, "type": "metricbeat", "uuid": "94b204f2-3dbb-4e68-9f67-18ba963e4ffc"}}}
2021-12-02T18:49:01.113+0545 INFO [beat] instance/beat.go:1023 Build info {"system_info": {"build": {"commit": "fd322dad6ceafec40c84df4d2a0694ea357d16cc", "libbeat": "7.15.2", "time": "2021-11-04T14:35:34.000Z", "version": "7.15.2"}}}
2021-12-02T18:49:01.114+0545 INFO [beat] instance/beat.go:1026 Go runtime info {"system_info": {"go": {"os": "linux", "arch": "amd64", "max_procs": 1, "version": "go1.16.6"}}}
2021-12-02T18:49:01.114+0545 INFO [beat] instance/beat.go:1030 Host info {"system_info": {"host": {"architecture": "x86_64", "boot_time": "2021-12-02T17:56:01+05:45", "containerized": false, "name": "samana-VM", "ip": ["127.0.0.1/8", "::1/128", "192.168.0.7/24", "fe80::d273:a838:40f6:44b3/64"]}, "kernel_version": "5

```

We can see the following result from metricbeat in kibana dashboard.





B. Collect metric from following sources in server1 and send them to elasticsearch. Store them in an index named "server1-metrics". a. Memory usage b. Disk usage c. Load average

To collect metrics we edit the config file as follows ( with load, memory and disk metrics sets)

```
GNU nano 4.8          metricbeat.yml          Modified
# Set to true to enable instrumentation of metricbeat.
#enabled: false

# Environment in which metricbeat is running on (eg: staging, production)
#environment: ""

# APM Server hosts to report instrumentation results to.
#hosts:
# - http://localhost:8200

# API Key for the APM Server(s).
# If api_key is set then secret_token will be ignored.
#api_key:

# Secret token for the APM Server(s).
#secret_token:

# ===== Migration =====>
# This allows to enable 6.7 migration aliases
#migration.6_to_7.enabled: true

metricbeat.modules:
- module: system
  metricsets:
    - load
    - memory
    - diskio
  enabled: true
  period: 5s
  index: "server1-metrics"
```

We can see that the collected metrics have been stored in an index names **server1-metrices**.

The screenshot shows the Elastic Stack Management interface. The left sidebar contains navigation links for Management, Ingest, Data, and Alerts and Insights. The main content area is titled 'Index Management' and includes tabs for Indices, Data Streams, Index Templates, and Component Templates. The 'Indices' tab is active, displaying a table of indices. The table has columns for Name, Health, Status, Primaries, Replicas, Docs count, Storage size, and Data stream. One index, 'server1-metrices', is listed with a yellow health status and an 'open' status. The interface also includes a search bar, filters for lifecycle status and phase, and a 'Reload indices' button.

**Management**

- Ingest ⓘ  
Ingest Node Pipelines
- Data ⓘ
  - Index Management**
  - Index Lifecycle Policies
  - Snapshot and Restore
  - Rollup Jobs
  - Transforms
  - Remote Clusters
- Alerts and Insights ⓘ  
Rules and Connectors

**Index Management** [Index Management docs](#)

**Indices** Data Streams Index Templates Component Templates

Update your Elasticsearch indices individually or in bulk. [Learn more.](#) [↗](#) ☐ Include rollout indices ☐ Include hidden indices

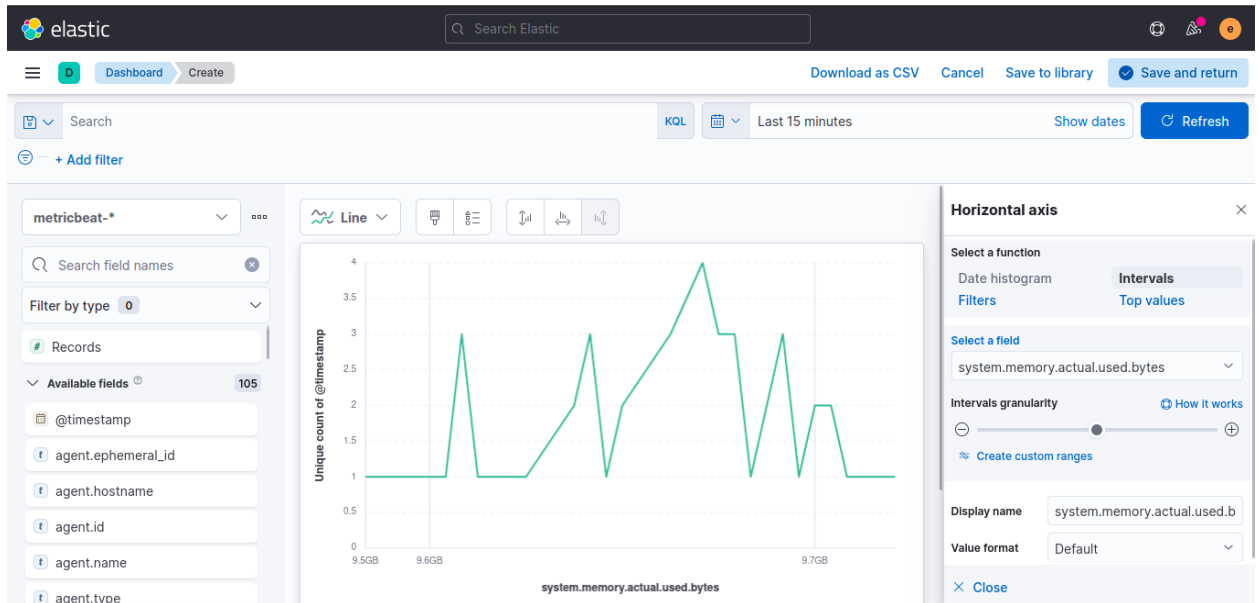
Lifecycle status Lifecycle phase [Reload indices](#)

<input type="checkbox"/>	Name	Health	Status	Primaries	Replicas	Docs count	Storage size	Data stream
<input type="checkbox"/>	<a href="#">server1-metrices</a>	● yellow	open	1	1	102	441.7kb	

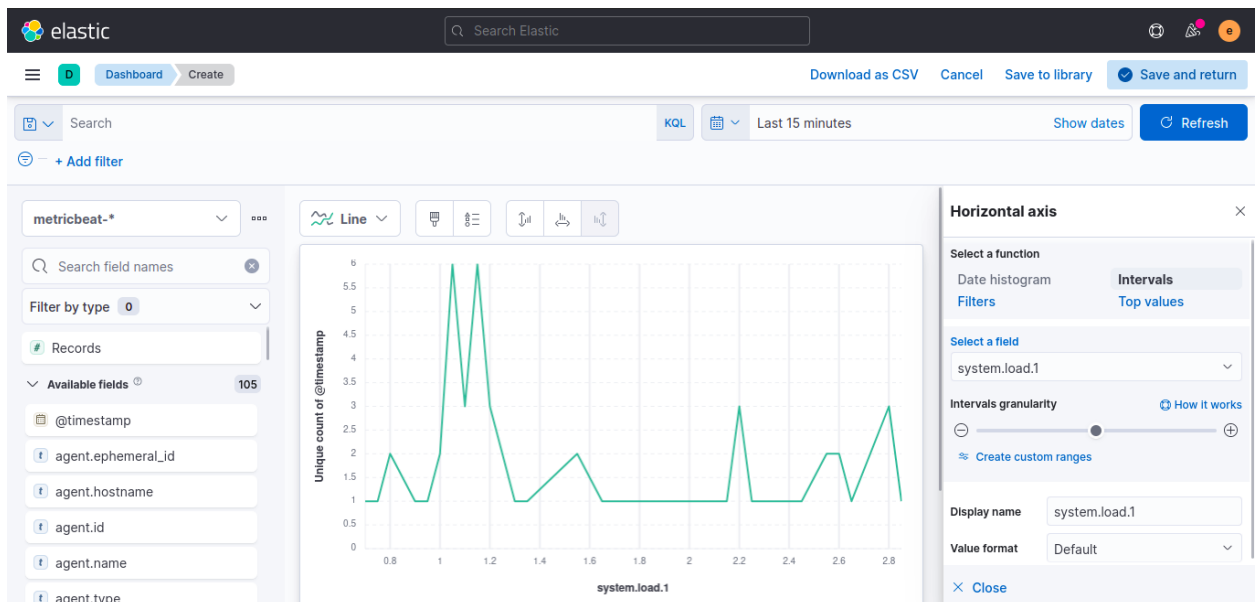
Rows per page: 10 [<](#) [1](#) [>](#)

1. Create a dashboard in kibana and generate visual report(line graph) for Memory usage and load average of server1 with relation to time

Line graph of memory usage with respect to time:

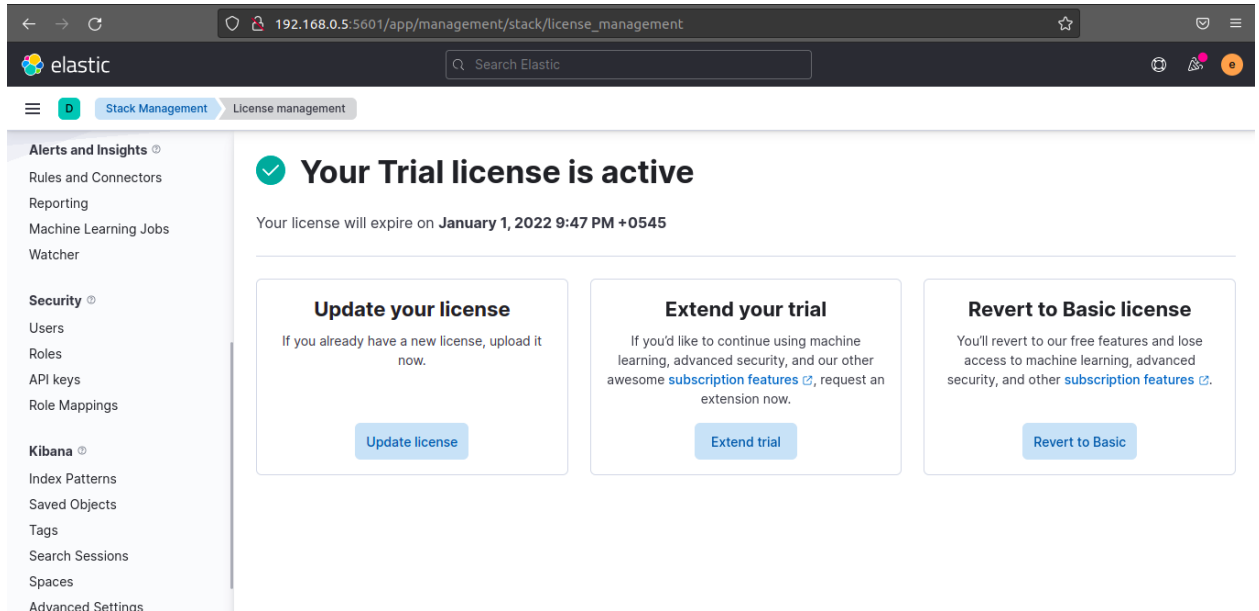


Line graph of load average in relation to time



2. **Generate alerts through kibana system for following thresholds**  
a. when memory usage > 80% for last 2 minutes send alert to a slack channel  
b. When Disk usage > 70% send alert to a slack channel  
c. When load average > 1 for last 2 minutes send alert to a slack channel

For creating alerts we start a trial license 30 day free plan.



We create rules as follows:

### a) For system load

Index patterns

metricbeat-7.15.2-2021.12.02-000001

Reset to default index patterns

Custom query

system.load.1 > 1

Import query from saved timeline

KQL

+ Add filter

Timeline template

None

Select which timeline to use when investigating generated alerts.

Quick query preview

Last hour

Preview results

Select a timeframe of data to preview query results

Continue

## b) For system memory

Index patterns

metricbeat-7.15.2-2021.12.02-000001 ×

Reset to default index patterns

Enter the pattern of Elasticsearch indices where you would like this rule to run. By default, these will include index patterns defined in Security Solution advanced settings.

Custom query

Import query from saved timeline

system.memory.actual.used.pct > 0.8

KQL

+ Add filter

Timeline template

None

Select which timeline to use when investigating generated alerts.

Quick query preview

Last hour

Preview results

Select a timeframe of data to preview query results

Continue

## c) For disk usage:

Index patterns

metricbeat-7.15.2-2021.12.02-000001 ×

Reset to default index patterns

Enter the pattern of Elasticsearch indices where you would like this rule to run. By default, these will include index patterns defined in Security Solution advanced settings.

Custom query

Import query from saved timeline

system.fsstat.total\_size.used > 70

KQL

+ Add filter

Timeline template

None

Select which timeline to use when investigating generated alerts.

Quick query preview

Last hour

Preview results

Select a timeframe of data to preview query results

We name the rules as follows:

2

About rule

Name

system load alert

Description

system.load is greater than 1

Default severity

Select a severity level for all alerts generated by this rule.

Low

2

About rule

Name

memory alert

Description

memory usage is greater than 80%

Default severity

Select a severity level for all alerts generated by this rule.

Low

☐ Severity override

About

Name

disk usage alert(>70%)

Description

.

We then schedule the rule(for all 3 cases) to run every 2 minutes:

3

Schedule rule

Runs every

2

Minutes

Rules run periodically and detect alerts within the specified time frame.

Additional look-back time

Optional

1

Minutes

Adds time to the look-back period to prevent missed alerts.

Continue

On the rule execution, we set the rule action to generate alert messages to slack channel using webhook integration.

4

Rule actions

Actions frequency

On each rule execution

Select when automated actions should be performed if a rule evaluates as true.

Actions

alert from system.load

Slack connector

Add connector




alert from system.load

Message

Rule {{context.rule.name}} generated {{state.signals\_count}} alerts



The created rules can be seen below:

Showing 3 rules	Selected 0 rules	Bulk actions	Refresh	Refresh settings					
<input type="checkbox"/>	Rule	Risk score	Severity	Last run	Last respo...	Last updated	Version	Tags	Activated
<input type="checkbox"/>	system load alert	77	Low	2 minutes ago	succeeded	Dec 3, 2021 @ 00:07:30.179	1	—	 ...
<input type="checkbox"/>	memory alert	58	Low	51 seconds ago	succeeded	Dec 3, 2021 @ 00:21:02.598	1	—	 ...
<input type="checkbox"/>	disk usage alert	21	Low	23 seconds ago	succeeded	Dec 3, 2021 @ 00:25:33.886	1	—	 ...

In order to trigger creating alerts, we perform the following action

```
samana@samana:~$ echo {1..1000000000}
```

Finally, we can see that in my slack channel all three kind of alerts are generated

**DevOps In...**

- Threads
- All DMs
- Mentions & reactions
- Slack Connect
- More

▼ Channels

- # class-recordings
- # devops
- # general
- # internship
- # random
- # team5
- + Add channels

▼ Direct messa...

- Slackbot
- Samana Pokhrel...

**Samana Pokhrel**

Message yourself? Why not! Think of this as a notepad – a place for jotting down a note or drawing up a to-do list.

[Edit your profile](#)

Today ▼

**Samana Pokhrel** 12:06 AM  
added an integration to this channel: [incoming-webhook](#)

**incoming-webhook** APP 12:38 AM  
Rule system load alert generated 12 alerts

**incoming-webhook** APP 12:44 AM  
Rule system load alert generated 13 alerts

**incoming-webhook** APP 8:14 AM  
Rule disk usage alert(>70%) generated 12 alerts  
Rule memory alert generated 11 alerts

Make a note of something

| **B** *I* ... Aa @ 😊 📎 ➤ ▼